Xutao Mao

xutao.henry.mao@gmail.com | (+86) 173-1925-3326 / (+852) 6767-4198
Website | GitHub | LinkedIn

Education

- **City University of Hong Kong** — *Hong Kong*
  Ph.D. Student, Computer Science — Sep 2026 – Aug 2030 (Expected)
  - Research Areas: AI Safety, Mechanistic Interpretability

- **Vanderbilt University** — *Nashville, US*
  Bachelor of Science, Honors in Computer Science & Mathematics — Aug 2022 – Dec 2025
  - Research Areas: LLM, Data Mining

Research Experience

- **City University of Hong Kong** — *Hong Kong*
  Graduate Researcher | Advisor: Dr. Cong Wang — Aug 2026 – Present
  - Acting and coordinating: how AI fails while they act in the world [AHA] [CAVE-Bench] [TrustFork]
  - Remembering and improving: how AI failures persist [PASB] [MisEvolve]
  - Looking inside and stepping in: how AI failures happens & how to mitigate failures [STARE] [Memory Interpret]

- **Hong Kong Baptist University** — *Hong Kong*
  Research Assistant | Advisor: Dr. Bo Han — Dec 2025 – Aug 2026
  - Focus: CoT Monitorability [TAME]

- **Vanderbilt University** — *Nashville, US*
  Undergraduate Researcher | Advisor: Dr. Tyler Derr — Feb 2025 – Dec 2025
  - Focus: LLM; Data Mining [Review Recommendation]

Publications

C=Conference, P=Preprint

[P.1] **Xutao Mao**, Rui Qian, Xiang Zheng‡, Cong Wang‡. **Agent Hacks Agents: Autoresearch Discovers Vulnerabilities in Production Agents**.

[P.2] **Xutao Mao***, Rui Qian*, Longxiang Wang, Xinjian Yi, Mingxuan Li, Linghan Chen, Yudong Gao, Xiang Zheng‡, Cong Wang‡. ***You're Right, Let Me Fix It*: How LLM Agents Damage Correct Work When Falsely Accused**.

[P.3] **Xutao Mao***, Rui Qian*, Linghan Chen, Yudong Gao, Junchi Liao, Jiulin Cai, Jinman Zhao, Cong Wang‡. **Trust the Brand, Lose Control: How Identity Hijacks LLM Agent Orchestration**.

[P.4] **Xutao Mao***, Liangjie Zhao*, Leyao Wang, Rui Qian, Qiang Huang, Wentao Wang, Bo Han‡, Xiang Zheng‡, Cong Wang‡. **Agents Don't Just Agree, They Remember: Benchmarking Persistent Sycophancy in Self-Improving Personal Agents**.

[P.5] **Xutao Mao**, Liangjie Zhao, Xiang Zheng‡, Cong Wang‡. **Practice Makes Unsafe: Skill Misevolution in Self-Improving LLM Agents**.

[P.6] Leyao Wang*, **Xutao Mao***, Xuhui Zhan, Yuying Zhao, Bo Ni, Ryan A. Rossi, Nesreen K. Ahmed, Tyler Derr‡. **Towards Bridging Review Sparsity in Recommendation with Textual Edge Graph Representation**.

[P.7] **Xutao Mao**, Jinman Zhao, Gerald Penn, Cong Wang‡. **What Happens Inside Agent Memory? Circuit Analysis from Emergence to Diagnosis**.

[C.1] **Xutao Mao***, Jianing Zhu*, Jinman Zhao, Tongliang Liu, Xiaowen Chu, Cong Wang‡, Bo Han‡. **Taming CoT Obfuscation in VLMs: From Mechanistic Evidence to Activation-Level Enforcement.** *NeurIPS 2026* **(Poster)**.

[C.2] Haobo Zhang*, **Xutao Mao***, Guangyuan Dong, Ziwei Li‡, Xuanbo Su, Kaijie Chen, Jing Yang, Zheng Lin. **MemMark: State-Evolution Attribution Watermarking for Agent Long-Term Memory Systems**. *EMNLP 2026* **(Finding)**.

[C.3] **Xutao Mao**, Liangjie Zhao, Tao Liu, Xiang Zheng‡, Hongying Zan, Cong Wang‡. **STARE: Step-wise Temporal Alignment and Red-teaming Engine for Multi-modal Toxicity Attack**. *ICML 2026* **(Poster)**.

[C.4] **Xutao Mao**‡, Ezra Tao, Leyao Wang. **MindVote: When AI Meets the Wild West of Social Media Opinion**. *AAAI 2026* **(Oral)**.

[C.5] Tao Liu†, **Xutao Mao**†, Hongying Zan‡, Dixuan Zhang, Yifan Li, Haixin Liu, Lulu Kong, Jiaming Hou, Rui Li, YunLong Li, Aoze Zheng, Zhiqiang Zhang, Zhewei Luo, Kunli Zhang, Min Peng. **LogicCat: Text-to-SQL Benchmark for Multi-Domain Reasoning Challenges**. *AAAI 2026* **(Poster)**.

Notes: *equal contribution; † alphabetical co-first/co-equal as noted; ‡corresponding author.

Work Experience

- **Pegasystems** *Cambridge, US*
  Software Development Engineer Intern May 2024 – Aug 2024

Honors and Awards

- **Hong Kong Postgraduate Scholarship** *2026*
  City University of Hong Kong

- **Vanderbilt Summer Research Program Scholarship** *2025*
  Vanderbilt University

Professional Service

- **Conference Reviewer:** ICLR '27; NeurIPS '26; AAAI '26 / '27; TheWebConf '26; ACL ARR (Oct 26)

Technical Skills & Interests

- **Models & Training:** LLM/VLM architectures (MoE, RoPE, FlashAttention); SFT, RL post-training (PPO, GRPO, DPO, RLHF), agentic RL, LoRA; PyTorch, DeepSpeed, FSDP, vLLM, verl; multi-node multi-GPU training
- **Research Interests:** AI Safety, Interpretability.